

EE2211 Tutorial 5

Dr Feng LIN

Input Output Given the following data pairs for training: $\{x = -10\} \rightarrow \{y = 5\}$ $\{x = -8\} \rightarrow \{y = 5\}$ $\{x = -3\} \rightarrow \{y = 4\}$ $\{x = -1\} \rightarrow \{y = 3\}$

Perform a linear regression with addition of a bias/offset term to the input feature vector and sketch the result of line fitting.

 $\{x = -2\} \rightarrow \{y = 2\}$ $\{x = -8\} \rightarrow \{y = 2\}$

- b) Perform a linear regression without inclusion of any bias/offset term and sketch the result of line fitting.
- What is the effect of adding a bias/offset term to the input feature vector?

(a) This is an over-determined system.

The input feature including bias/offset can be written as

$$X^{T} = \begin{bmatrix} 1 & 1 & 1 & 1 & 1 & 1 \\ -10 & -8 & -3 & -1 & 2 & 8 \end{bmatrix}$$

$$\widehat{w} = (X^T X)^{-1} X^T y = \begin{bmatrix} 6 & -12 \\ -12 & 242 \end{bmatrix} \begin{bmatrix} 1 & 1 & 1 & 1 & 1 \\ -10 & -8 & -3 & -1 & 2 & 8 \end{bmatrix} \begin{bmatrix} 3 \\ 5 \\ 4 \\ 3 \\ 2 \\ 2 \end{bmatrix} = \begin{bmatrix} 3.1055 \\ -0.1972 \end{bmatrix}$$

(b) This is an over-determined system.

In this case, the input feature without inclusion of bias/offset is a vector given by $\begin{bmatrix} -10 & -8 & -3 & -1 & 2 & 8 \end{bmatrix}^T$.

$$\widehat{w} = (x^T x)^{-1} x^T y = [242]^{-1} [-10 \quad -8 \quad -3 \quad -1 \quad 2 \quad 8] \begin{bmatrix} 5\\4\\3\\2\\2 \end{bmatrix} = -0.3512$$

(c) The bias/offset term allows the line to move away from the origin (moved vertically in this case).

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from numpy.linalg import inv

# define the matrix X, bias vector and y

X = np.array([[-10], [-8], [-3], [-1], [2], [8]])
b = np.ones( (len(X),1) )

X_b = np.hstack((b, X)) # X matrix with bias

y = np.array([[5], [5], [4], [3], [2], [2]])
```

#(a) Perform a linear regression with addition of a bias/offset term
w_b = inv(X_b.T@X_b)@X_b.T@y

```
#(b) Perform a linear regression without inclusion of any bias/offset term
W = inv(X.T@X)@X.T@y
                                   Left-inverse
print("w=", w)
from numpy.linalg import pinv
                                   pseudoinverse
w p = pinv(X)@y
print("w p=", w p)
                                    Cutoff condition
                                       number
from numpy.linalg import lstsq
                                                       Least square solution
w_ls, residuals, rank, s = lstsq(X, y, rcond=-1)
print("w_sl", w_ls)
W = [[-0.35123967]]
w_p = [[-0.35123967]]
w_sl [[-0.35123967]]
```

© Copyright National University of Singapore. All Rights Reserved.

```
# show the effect of adding a bias/offset term
X \text{ test} = \text{np.linspace}(-10,10, 400)
X_test = X_test.reshape(-1, 1)
b test = np.ones((len(X test),1)) # generate a bias column vector
X_b_test = np.hstack((b_test, X test))
y b test = X b test@w b
plt.figure()
plt.plot(X test, y b test, color='red', label='Linear Regression')
plt.plot(X, y, 'o', label='Training Samples')
plt.xlim(-10,10)
plt.ylim(-6, 6)
plt.grid(True)
plt.legend()
plt.show()
```

(c) The bias/offset term allows the line to move away from the origin (moved vertically in this case).

Linear Regression

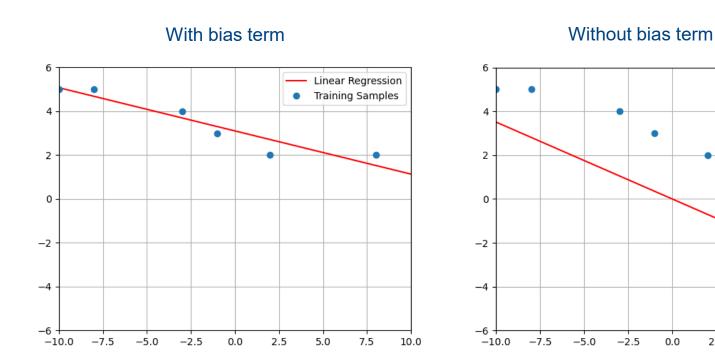
Training Samples

2.5

5.0

7.5

10.0





(Linear Regression, prediction, even/under-determined)

Given the following data pairs for training:

$$\{x_1 = 1, \quad x_2 = 0, \quad x_3 = 1\} \rightarrow \{y = 1\}$$

 $\{x_1 = 2, \quad x_2 = -1, \quad x_3 = 1\} \rightarrow \{y = 2\}$
 $\{x_1 = 1, \quad x_2 = 1, \quad x_3 = 5\} \rightarrow \{y = 3\}$

- (a) Predict the following test data without inclusion of an input bias/offset term.
- (b) Predict the following test data with inclusion of an input bias/offset term.

$$\{x_1 = -1, \quad x_2 = 2, \quad x_3 = 8\} \rightarrow \{y = ?\}$$

 $\{x_1 = 1, \quad x_2 = 5, \quad x_3 = -1\} \rightarrow \{y = ?\}$

(a) Without bias, this is an even-determined system and *X* is invertible

$$\widehat{w} = X^{-1}y = \begin{bmatrix} 1 & 0 & 1 \\ 2 & -1 & 1 \\ 1 & 1 & 5 \end{bmatrix}^{-1} \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix} = \begin{bmatrix} 0.3333 \\ -0.6667 \\ 0.6667 \end{bmatrix}$$

$$\hat{\mathbf{y}}_{t} = X_{t}^{-1} \hat{\mathbf{w}} = \begin{bmatrix} -1 & 2 & 8 \\ 1 & 5 & -1 \end{bmatrix} \begin{bmatrix} 0.3333 \\ -0.6667 \\ 0.6667 \end{bmatrix} = \begin{bmatrix} 3.6667 \\ -3.6667 \end{bmatrix}$$

(b) After adding bias, it becomes an under-determined system.

$$\widehat{w} = X^{\mathrm{T}}(XX^{\mathrm{T}})^{-1}y = \begin{bmatrix} 1 & 1 & 1 \\ 1 & 2 & 1 \\ 0 & -1 & 1 \\ 1 & 1 & 5 \end{bmatrix} \begin{bmatrix} 3 & 4 & 7 \\ 4 & 7 & 7 \\ 7 & 7 & 28 \end{bmatrix}^{-1} \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix} = \begin{bmatrix} -0.1429 \\ 0.5238 \\ -0.4762 \\ 0.6190 \end{bmatrix}$$

$$\hat{\mathbf{y}}_{t} = X_{t} \ \hat{\mathbf{w}} = \begin{bmatrix} 1 & -1 & 2 & 8 \\ 1 & 1 & 5 & -1 \end{bmatrix} \begin{bmatrix} -0.1429 \\ 0.5238 \\ -0.4762 \\ 0.6190 \end{bmatrix} = \begin{bmatrix} 3.3333 \\ -2.6190 \end{bmatrix}$$

```
In [3]: import numpy as np
        import pandas as pd
        import matplotlib.pyplot as plt
        from numpy.linalg import inv
        # Without bias, this is an even-determined system and X is invertible.
            = np.array([[1, 0, 1], [2, -1, 1], [1, 1, 5]])
        y = np.array([[1], [2], [3]])
        b = np.ones((len(X),1))
        X b = np.hstack((b, X)) # X matrix with bias
In [4]: #(a) Perform a linear regression with addition of a bias/offset term
        w = inv(X)@y
        print(w)
        X t = np.array([[-1, 2, 8], [1, 5, -1]])
        v t = X t_0 w
        print(y t)
        [[ 0.33333333]
         [-0.66666667]
         [ 0.66666667]]
        [[ 3.66666667]
         [-3.66666667]]
```

In [5]: #(b) After adding bias, it becomes an under-determined system.
w_b = X_b.T@inv(X_b@X_b.T)@y

(Linear Regression, prediction, extrapolation)

A college bookstore must order books two months before each semester starts. They believe that the number of books that will ultimately be sold for any particular course is related to the number of students registered for the course when the books are ordered.

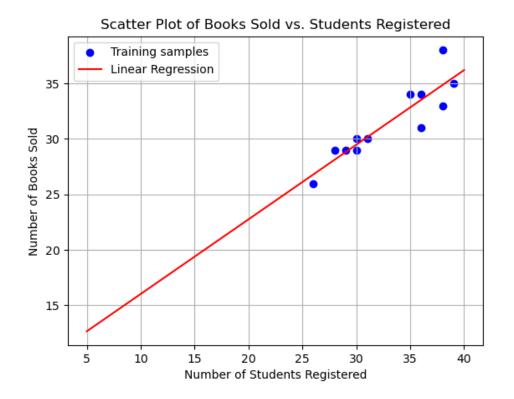
They would like to develop a linear regression equation to help plan how many books to order. From past records, the bookstore obtains the number of students registered, X, and the number of books actually sold for a course, Y, for 12 different semesters. These data are shown below.

Semester	Students	Books
1	36	31
2	28	29
3	35	34
4	39	35
5	30	29
6	30	30
7	31	30
8	38	38
9	36	34
10	38	33
11	29	29
12	26	26

- (a) Obtain a scatter plot of the number of books sold versus the number of registered students.
- (b) Write down the regression equation and calculate the coefficients for this fitting.
- (c) Predict the number of books that would be sold in a semester when 30 students have registered.
- (d) Predict the number of books that would be sold in a semester when 5 students have registered.

Question 3

(a) Scatter plot



Regression equation: y = Xw

(c)

$$\widehat{y}_t = X_t \widehat{w} = \begin{bmatrix} 1 & 30 \end{bmatrix} \begin{bmatrix} 9.30 \\ 0.6727 \end{bmatrix} = 29.4818$$

(d) $(\hat{y}_t = 12.6636)$ This prediction appears to be somewhat over optimistic. Since 5 students is not within the range of the sampled number of students, it might not be appropriate to use the regression equation to make this prediction. We do not know if the straight-line model would fit data at this point, and we might not want to extrapolate far beyond the observed range.

```
import numpy as np
import matplotlib.pyplot as plt
from scipy import stats
from numpy.linalg import inv
# Define Data: Number of students (X) and number of books sold (v)
X = \text{np.array}([36, 28, 35, 39, 30, 30, 31, 38, 36, 38, 29, 26])
X = X.reshape(-1, 1) # reshape as a vertical vector
y = np.array([31, 29, 34, 35, 29, 30, 30, 38, 34, 33, 29, 26])
y = y.reshape(-1, 1) # reshape as a vertical vector
b = np.ones((len(X), 1))
X b = np.hstack((b, X)) # add bias to the X matrix
# (a) Scatter plot
plt.scatter(X, y, color='blue', label='Training samples')
plt.title('Scatter Plot of Books Sold vs. Students Registered')
plt.xlabel('Number of Students Registered')
plt.vlabel('Number of Books Sold')
plt.grid(True)
plt.legend()
```

```
# (b) Linear Regression: Calculate w
W = inv(X_b.T_0X_b)_0X_b.T_0y
print("w is"); print(w)
 # draw the estimated Line
 X_t = np.linspace(5,40,50)
X t = X t.reshape(-1, 1)
b_t = np.ones((len(x_t),1)) # generate a bias column vector
X b t = np.hstack((b t, X t))
v b t = X_b_t@w
 plt.plot(X_t, y_b_t, color='red', label='Linear Regression')
 plt.legend(); plt.show()
 # (c) Predict books sold when 30 students are reaistered
X \text{ new } 30 = 30
 y_pred_30 = np.array([ [1, X_new_30]]) @ w
 print(y pred 30)
 # (d) Predict books sold when 5 students are registered
 X new 5 = 5
 y_pred_5 = np.array([ [1, X_new_5]]) @ w
 print(y pred 5)
```

Repeat the above problem using the following training data:

- a) Calculate the regression coefficients for this fitting.
- b) Predict the number of books that would be sold in a semester when 30 students have registered.
- c) Purge those duplicating data and re-fit the line and observe the impact on predicting the number of books that would be sold in a semester when 30 students have registered.
- d) Sketch and compare the two fitting lines.

	<u> </u>
Students	Books
36	31
26	20
35	34
39	35
26	20
30	30
31	30
38	38
36	34
38	33
26	20
26	20
	Students 36 26 35 39 26 30 31 38 36 38 26

Using the full data:

$$\widehat{y}_t = X_t \widehat{w} = \begin{bmatrix} 1 & 30 \end{bmatrix} \begin{bmatrix} -10.4126 \\ 1.2143 \end{bmatrix} = 26.0177$$

After having the duplicating data purged:
$$\widehat{w} = (X^T X)^{-1} X^T y = \begin{bmatrix} 9 & 309 \\ 309 & 10763 \end{bmatrix}^{-1} \begin{bmatrix} 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 36 & 35 & 39 & 30 & 31 & 38 & 36 & 38 & 26 \end{bmatrix} \begin{bmatrix} 31 \\ 34 \\ 35 \\ 30 \\ 30 \\ 38 \\ 34 \\ 33 \\ 20 \end{bmatrix} = \begin{bmatrix} -3.5584 \\ 1.0260 \end{bmatrix}$$

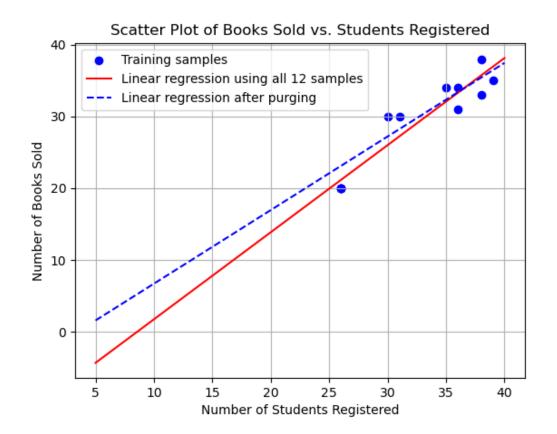
$$\widehat{y_t} = X_t \widehat{w} = \begin{bmatrix} 1 & 30 \end{bmatrix} \begin{bmatrix} -3.5584 \\ 1 & 0260 \end{bmatrix} = 27.2208$$

Note: these results show that duplicating samples can influence the learning and decision too. In this case, purging seems to give a more optimistic prediction for a relatively small number of students (< 37) and more conservative prediction for a relatively large number of students (>37).

```
import numpy as np
import matplotlib.pyplot as plt
from scipy import stats
from numpy.linalg import inv
# Define Data: Number of students (X) and number of books sold (y)
X = np.array([36, 26, 35, 39, 26, 30, 31, 38, 36, 38, 26, 26])
X = X.reshape(-1, 1) # reshape as a vertical vector
y = np.array([31, 20, 34, 35, 20, 30, 30, 38, 34, 33, 20, 20])
y = y.reshape(-1, 1) # reshape as a vertical vector
b = np.ones((len(X), 1))
X b = np.hstack((b, X)) # add bias to the X matrix
# Scatter plot
plt.scatter(X, y, color='blue', label='Training samples')
plt.title('Scatter Plot of Books Sold vs. Students Registered')
plt.xlabel('Number of Students Registered')
plt.ylabel('Number of Books Sold')
plt.grid(True)
plt.legend()
# (a) Linear Regression: Calculate w
w = inv(X b.T@X b)@X b.T@y
print("w is"); print(w)
# draw the estimated line
X t = np.linspace(5,40,50)
X t = X t.reshape(-1, 1)
b t = np.ones((len(X t),1)) # generate a bias column vector
X_b_t = np.hstack((b_t, X_t))
y t = X b t@w
plt.plot(X t, y t, color='red', label='Linear regression using all 12 samples')
```

```
\# X \text{ cleaned = np.array}([36, 35, 39, 30, 31, 38, 36, 38, 26])
\# X \text{ cleaned} = X \text{ cleaned.reshape}(-1,1)
# y cleaned = np.array([31, 34, 35, 30, 30, 38, 34, 33, 20])
\# y cleaned = y cleaned.reshape(-1,1)
# find the unique data
duplicated = np.hstack((X,v))
cleaned = np.unique(duplicated , axis=0)
print("cleaned data = ")
print(cleaned)
X cleaned = cleaned[:,0]
X cleaned = X cleaned.reshape(-1,1)
y_cleaned = cleaned[:,1]
y cleaned = y cleaned.reshape(-1,1)
b cleaned = np.ones( (len(X cleaned), 1) )
X_b_cleaned = np.hstack((b_cleaned, X_cleaned)) # add bias to the X matrix
print(X b cleaned)
w cleaned = inv(X b cleaned.T@X b cleaned)@X b cleaned.T@y cleaned
print(w cleaned)
# (d) Sketch and compare the two fitting lines
y t cleaned = X b t@w cleaned
plt.plot(X t, y t cleaned, color='blue', linestyle='--', label='Linear regression after purging')
plt.legend()
plt.show()
```

(c) Purge duplicates (keep only one instance where X = 26 and Y = 20)



Download the data file "government-expenditure-on-education.csv" from Canvas Tutorial Folder. It depicts the government's educational expenditure over the years (downloaded in July 2021 from https://data.gov.sg/dataset/government-expenditure-on-education)

Predict the educational expenditure of year 2021 based on linear regression. Solve the problem using Python with a plot. Note: please use the file from the canvas link.

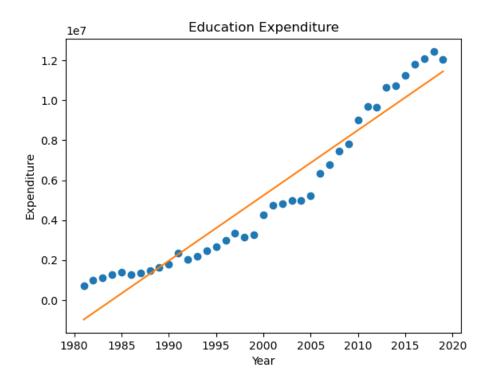
Hint: use Python packages like numpy, pandas, matplotlib.pyplot, numpy.linalg.

df.head()

	year	recurrent_expenditure_total
0	1981	712732
1	1982	983751
2	1983	1107113
3	1984	1272559
4	1985	1388186

import numpy as np

```
import pandas as pd
import matplotlib.pyplot as plt
from numpy.linalg import inv
# read data
df = pd.read_csv("government-expenditure-on-education.csv")
# convert the data to matrices: X and y
expenditureList = df ['recurrent expenditure total'].tolist()
yearList = df ['year'].tolist()
m_list = [[1]*len(yearList), yearList]
X = np.array(m_list).T
y = np.array(expenditureList)
# linear regression
w = inv(X.T @ X) @ X.T @ y
print(w)
# plot the results
v line = X.dot(w)
plt.plot(yearList, expenditureList, 'o', label = 'Expenditure over the years')
plt.plot(yearList, y_line)
plt.xlabel('Year')
plt.ylabel('Expenditure')
plt.title('Education Expenditure')
plt.show()
# prediction
y_predict = np.array([1, 2021]).dot(w)
print(y predict)
```



Answer: The predicted educational expenditure in year 2021 is 12102904.270637512

Download the CSV file for red-wine using "wine = pd.read_csv("https://archive.ics.uci.edu/ml/machine-learningdatabases/wine-quality/winequality-red.csv",sep=';') ". Use Python to perform the following tasks.

Hint: use Python packages like numpy, pandas, matplotlib.pyplot, numpy.linalg, and sklearn.metrics.

- a) Take y = wine.quality as the target output and x = wine.drop('quality',axis = 1) as the input features. Assume the given list of data is already randomly indexed (i.e., not in particular order), split the database into two sets: [0:1500] samples for regression training, and [1500:1599] samples for testing.
- b) Perform linear regression on the training set and print out the learned parameters.
- c) Perform prediction using the test set and provide the prediction accuracy in terms of the mean of squared errors (MSE).

<class 'pandas.core.frame.DataFrame'> RangeIndex: 1599 entries, 0 to 1598 Data columns (total 12 columns): Column Non-Null Count Dtype fixed acidity 1599 non-null float64 volatile acidity 1599 non-null float64 citric acid 1599 non-null float64 residual sugar 1599 non-null float64 4 chlorides 1599 non-null float64 5 free sulfur dioxide 1599 non-null float64 6 total sulfur dioxide 1599 non-null float64 density 1599 non-null float64 pН 1599 non-null float64 sulphates 1599 non-null float64 10 alcohol 1599 non-null float64 11 quality 1599 non-null int64 dtypes: float64(11), int64(1)

[4]: wine

	C		-14-114			6	4-4-1	4	-11		-11-1	
	тіхед асідіту	volatile acidity	citric acid	residuai sugar	cniorides	Tree suitur dioxide	total sulfur dioxide	density	рн	suipnates	alconol	quality
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	5
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68	9.8	5
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65	9.8	5
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	9.8	6
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	5
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	10.5	5
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	11.2	6
1596	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75	11.0	6
1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	10.2	5
1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66	11.0	6

import pandas as pd

```
#import matplotlib.pyplot as plt
              import numpy as np
             from numpy.linalg import inv
              from sklearn.metrics import mean_squared_error
              ## get data from web
              # wine = pd.read csv("https://archive.ics.uci.edu/ml/machine-learning-databases/winequality/winequality-red.csv",sep=';')
             wine = pd.read csv("winequality-red.csv", sep=';') # get data from the downloaded local file
             # wine.info()
             y = wine.quality

    x = wine.drop('quality', axis=1): creates a variable x that stores the features by dropping the

             x = wine.drop('quality',axis = 1)
                                                     quality column from 'wine';

    The axis=1 argument indicates that the column (quality) is being removed (axis 1 is for columns,

             ## Include the offset/bias term
             x0 = np.ones((len(y),1))
                                                     axis 0 is for rows).
             X = np.hstack((x0,x))
             ## split data into training and test sets
             ## (Note: this exercise introduces the basic protocol of using the training-test
              ## partitioning of samples for evaluation assuming the list of data is already randomly indexed)
              ## In case you really want a general random split to have a better training/test distributions:
              ## from sklearn.model_selection import train_test_split
              ## train X, test X, train y, test y = train test split(X, y, test size=99/1599, random state = 0)
              # randomly split the data into training and testing sets
                                                                                                    • test size=99/1599: This specifies the proportion of
             from sklearn.model selection import train test split
                                                                                                       the dataset to include in the test set
              train_X,test_X,train_y,test_y = train_test_split(X,y,test_size=99/1599, random state = 4)
                                                                                                      random state=4: This ensures reproducibility.
             # deterministically split data into training and testing sets
             # train X = X[0:1500]
                                           Split the database into two sets: [0:1500] samples for regression training,
             # train y = y[0:1500]
             # test X = X[1500:1599]
                                           and [1500:1599] samples for testing.
             # test y = y[1500:1599]
             ## linear regression
             w = inv(train_X.T @ train_X) @ train_X.T @ train_y
             print(w)
             yt est = test X.dot(w);
             MSE = np.square(np.subtract(test y,yt est)).mean()
             print(MSE)
© Copyright National I
```

MSE = mean squared error(test y,yt est) print(MSE)

Results:

```
[ 2.22330327e+01 2.68702621e-02 -1.12838019e+00 -2.06141685e-01
```

1.22000584e-02 -1.77718503e+00 4.29357454e-03 -3.18953315e-03

-1.81795124e+01 -3.98142390e-01 8.92474793e-01 2.77147239e-01]

- 0.34352638121356655
- 0.34352638121356655

This question is related to understanding of modelling assumptions. The function given by $f(\mathbf{x}) = 1 + x_1 + x_2 - x_3 - x_4$ is affine.

- a) True
- b) False

MCQ: There could be more than one answer.

Suppose $f(\mathbf{x})$ is a scalar function of d variables where \mathbf{x} is a $d \times 1$ vector. Then, without taking data points into consideration, differentiation of $f(\mathbf{x})$ $w.r.t.\mathbf{x}$ is

- a) a scalar
- b) a $d \times 1$ vector
- c) a $d \times d$ matrix
- d) a $d \times d \times d$ tensor
- e) None of the above

(Linear regression with multiple outputs)

The values of feature vector **x** and their corresponding values of target vector **y** are shown in the table below:

x	[3, -1, 0]	[5, 1, 2]	[9, -1, 3]	[-6, 7, 2]	[3, -2, 0]
y	[1, -1]	[-1, 0]	[1, 2]	[0, 3]	[1, -2]

Find the least square solution of w using linear regression of multiple outputs and then estimate the value of y when x = [8, 0, 2].

```
#python
import numpy as np
from numpy.linalg import inv
X = \text{np.array}([[1, 3, -1, 0], [1, 5, 1, 2], [1, 9, -1, 3], [1, -6, 7, 2], [1, 3, -2, 0]])
Y = np.array([[1, -1], [-1, 0], [1, 2], [0, 3], [1, -2]])
W = inv(X.T @ X) @ X.T @ Y
print(W)
newX=np.array([1, 8, 0, 2])
newY=newX@W
print(newY)
[[ 1.14668974 -0.95997404]
 [-0.630463 -0.33427088]
 [-1.10601471 -0.24426655]
```

[-1.17784509 -0.07507572]

THANK YOU